

Sources of Variability in Solutions to Getting and Cleaning Data Coursera Project

ANDREW LEROUX¹

¹*Johns Hopkins University, Baltimore, MD 21205, USA*

October 11, 2017

1 Introduction

In any data analysis, the first step is to first collect the data. Once collected, there is usually a data processing/cleaning procedure that takes the raw data and transforms it into an “analysis ready” format using some programming language. How individuals navigate this procedure is of scientific interest for several reasons. Among those reasons are: (1) the importance of upstream data processing pipelines to downstream analysis – i.e. garbage in, garbage out; and (2) how heterogeneity in individuals’ coding styles results in readability and understandability among their peers – i.e. “can this code reasonably be maintained by someone else”.

While a comprehensive answer to the question “how do people code?” is beyond the scope of this paper, we attempt to identify the key features which can explain the variability in the way people solve a particular coding problem: processing data. Analyzing how people process data has the potential to address both of the aforementioned scientific questions.

We use publicly available data from students who enrolled in the Coursera class “Getting and Cleaning Data” (include citation). The final project for the class has students merge two data sets, rename variables, and then summarize the data in a smaller, tidy dataset. To perform this data processing, students used the statistical software *R* (include citation), but were otherwise given complete flexibility on how to complete the assignment. In addition to students’ code, we have a search ranking score provided by Github, which we use as a proxy for a “rating” or “grade” on students’ assignments. We associate with score with

The rest of this document is organized as follows: In section 2 we describe the data retrieval procedure and limitations of this procedure in detail, as well as the statistical methods used in data analysis. In section 3 we present the results of our statistical analyses. We conclude with a high level discussion of our results and implications for future research in section 4.

2 Methods

Here we discuss our data collection procedure and the statistical methods used in our data analysis.

2.1 Getting the Data

The data was retrieved from Github. The data was retrieved over the 72 hour period period 10/02/2017-10/04/2017. All data retrieval and analyses were performed in *R* (add package + R citation). At a high level, the retrieval procedure followed the two steps below

- Use the *gh* package available (add package citation) to find all repos associated with the class which were created on or after 12/01/2007
- Use the *gh* package to search the repos discovered in the previous step for a file called “run_analysis.R”. Then scrape the data using the URL structure implied by the location of the “run_analysis.R” file using the *readlines* function

More detailed information on each of these steps is provided below.

2.1.1 Finding Repo Names

Searching for repos associated with the getting and cleaning data class results in over 30,000 search results. However, the github API will only report up to 1,000 search results. To get around this, we searched repos by date of creation considering periods of two weeks. The creation dates examined spanned the dates 12/01/2007-10/02/2017. The justification for the start date was based on the start of the creation date of the class as well as some manual exploration which suggested the search results turned up zero entries for reasonable creation dates beyond this period. Note that using a 2 week moving window we never hit the 1,000 search result limit, implying that our window was sufficiently high resolution for this data. In another application (or for the class going forward), this window may need to shrink in order to capture all repos created during the search window.

The exact repo search was performed using the query: “GET /search/repositories?q=getting+and+cleaning+data+repo_date&per_page=100” where **repo_date** is of the form “created:2007-11-31..2007-12-15” to get repos created anytime during the two week period 12/01/2007-12/14/2007.

We note that one limitation of our procedure is that any repos created prior to 12/01/2007 would not be included in our analysis. However, since this date is prior to the creation of the class, in order for someone to be missed using this procedure, they would've had to rename or re-purpose an existing repo. We believe this to be unlikely, but note that it is a possibility. Also, note that our code does not allow for repository names to have hyphens in them. These individuals were excluded from our analysis. These individuals number **INCLUDE NUMBER**.

2.1.2 Scraping the Data

To scrape the data, we looped over the repos found in the previous step and searched their (entire) repository for a file called "run_analysis.R". To do so we used the following search query in the **gh()** function: "GET /search/code?q=repo:**repo_name**+extension:r" where **repo_name** is the repository name. This will search the entire repository for any .R files.

We then used regular expressions (*gregexpr*) to find whether any of these .R files matched "run_analysis.r" by using the *to.lower* function in R to allow for various capitalized letters and still match. Finally, we scraped the "run_analysis.R" file using the *readlines* function on the appropriate URL based on the file name, repository name and branch name.

We note that our code does impose some limitations. First and foremost, there were some files that we attempted to read, but were found to be non- UTF8 files. We did not try to parse these files at all and they were excluded from the analysis. From manual inspection, these tended to be individuals who copied and pasted R console output into a .R script and used that as their final project. So these are users who, sometimes, were able to complete the task. We estimate the number of these individuals to be **INCLUDE NUMBER**.

2.1.3 Limitations of the Data Retrieval Procedure

Discuss limitations of both my code for extracting unique functions (i.e. error rate) as well as the limitations of the Github "score"

2.2 Analyzing the Data

Discuss choice and justification for the 10 features explored.

2.2.1 Principal components analysis

Discuss PCA.

2.2.2 Associating Github “scores” with features of the data

Discuss linear regression of github scores on features/PC scores.

3 Results

3.1 Principal Components Analysis

Feature	PC 1	PC 2	PC 3	PC 4	PC 5	PC 6
# of lines of code	-0.47	0.15	0.02	-0.18	0.15	-0.07
# of lines of comments	-0.35	0.11	-0.27	-0.44	0.21	0.29
# of blank lines	-0.36	0.15	0.06	-0.35	0.2	-0.51
# of characters of code	-0.34	0.31	0.28	0.51	0.02	-0.04
# of packages	-0.1	-0.03	-0.66	0.28	-0.3	-0.57
# of object assignments	-0.31	-0.54	0.21	0.04	0.07	-0.17
# of unique assignment names	-0.2	-0.68	0.13	0.14	0.08	0.01
# of base subsetting used	-0.19	0.01	0.29	-0.32	-0.87	0.02
# of functions called	-0.39	0.23	0.08	0.44	-0.06	0.24
# of unique functions called	-0.28	-0.19	-0.5	0.04	-0.14	0.49
% Variance explained	38.3	15.2	12.7	10.5	8.3	6.8
Cumulative % variance explained	38.3	53.5	66.2	76.7	85	91.9

Table 1: *First Six Principal Components*

3.2 Regression Analysis

3.3 Associating Github Score with Code Features

4 Discussion

Ultimately, we found that features which closely correspond to our perceptions of what “efficient” code looks like explain the majority of the variability in how individuals

solved this particular assignment and are most predictive of the Github search score.